



Want help with deep learning for text? [Take the FREE Mini-Course](#)



How to Use Small Experiments to Develop a Caption Generation Model in Keras

by **Jason Brownlee** on November 24, 2017 in **Deep Learning for Natural Language Processing**

Caption generation is a challenging artificial intelligence problem where a textual description must be generated for a photograph.

It requires both methods from computer vision to understand the content of the image and a language model from the field of natural language processing to turn the understanding of the image into words in the right order. Recently, deep learning methods have achieved state of the art results on examples of this problem.

It can be hard to develop caption generating models on your own data, primarily because the datasets and the models are so large and take days to train. An alternative approach is to explore model configurations with a small sample of the fuller dataset.

In this tutorial, you will discover how you can use a small sample of a standard photo captioning dataset to explore different deep model designs.

After completing this tutorial, you will know:

- How to prepare data for photo captioning modeling.
- How to design a baseline and test harness to evaluate the skill of models and control for their stochastic nature.
- How to evaluate properties like model skill, feature extraction models, and word embeddings in order to lift model skill.

Let's get started.



How to Use Small Experiments to Develop a Caption Generation Model in Keras
Photo by [Per](#), some rights reserved.

Tutorial Overview

This tutorial is divided into 6 parts; they are:

1. Data Preparation
2. Baseline Caption Generation Model
3. Network Size Parameters
4. Configuring the Feature Extraction Model
5. Word Embedding Models
6. Analysis of Results

Python Environment

This tutorial assumes you have a Python SciPy environment installed, ideally with Python 3.

You must have Keras (2.0 or higher) installed with either the TensorFlow or Theano backend.

The tutorial also assumes you have scikit-learn, Pandas, NumPy, and Matplotlib installed.

If you need help with your environment, see this tutorial:

- [How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda](#)

I recommend running the code on a system with a GPU.

You can access GPUs cheaply on Amazon Web Services. Learn how in this tutorial:

- [How To Develop and Evaluate Large Deep Learning Models with Keras on Amazon Web Services](#)

Let's dive in.

Need help with Deep Learning for Text Data?

Take my free 7-day email crash course now (with code).

Click to sign-up and also get a free PDF Ebook version of the course.

[Start Your FREE Crash-Course Now](#)

Data Preparation

First, we need to prepare the dataset for training the model.

We will use the Flickr8K dataset that is comprised of a little more than 8,000 photographs and their descriptions.

You can download the dataset from here:

- [Framing image description as a ranking task: data, models and evaluation metrics.](#)

Unzip the photographs and descriptions into your current working directory into *Flicker8k_Dataset* and *Flicker8k_text* directories respectively.

There are two parts to the data preparation, they are:

1. Preparing the Text
2. Preparing the Photos

Preparing the Text

The dataset contains multiple descriptions for each photograph and the text of the descriptions requires some minimal cleaning.

First, we will load the file containing all of the descriptions.

```

1 # load doc into memory
2 def load_doc(filename):
3     # open the file as read only
4     file = open(filename, 'r')
5     # read all text
6     text = file.read()
7     # close the file
8     file.close()
9     return text
10
11 filename = 'Flickr8k_text/Flickr8k.token.txt'
12 # load descriptions
13 doc = load_doc(filename)

```

Each photo has a unique identifier. This is used in the photo filename and in the text file of descriptions. Next, we will step through the list of photo descriptions and save the first description for each photo. Below defines a function named *load_descriptions()* that, given the loaded document text, will return a dictionary of photo identifiers to descriptions.

```

1 # extract descriptions for images
2 def load_descriptions(doc):
3     mapping = dict()
4     # process lines
5     for line in doc.split('\n'):
6         # split line by white space
7         tokens = line.split()
8         if len(line) < 2:
9             continue
10        # take the first token as the image id, the rest as the description
11        image_id, image_desc = tokens[0], tokens[1:]
12        # remove filename from image id
13        image_id = image_id.split('.')[0]
14        # convert description tokens back to string
15        image_desc = ' '.join(image_desc)
16        # store the first description for each image
17        if image_id not in mapping:
18            mapping[image_id] = image_desc
19    return mapping
20
21 # parse descriptions
22 descriptions = load_descriptions(doc)
23 print('Loaded: %d ' % len(descriptions))

```

Next, we need to clean the description text.

The descriptions are already tokenized and easy to work with. We will clean the text in the following ways in order to reduce the size of the vocabulary of words we will need to work with:

- Convert all words to lowercase.
- Remove all punctuation.
- Remove all words that are one character or less in length (e.g. 'a').

Below defines the *clean_descriptions()* function that, given the dictionary of image identifiers to descriptions, steps through each description and cleans the text.

```

1 import string
2
3 def clean_descriptions(descriptions):
4     # prepare translation table for removing punctuation

```

```

5     table = str.maketrans('', '', string.punctuation)
6     for key, desc in descriptions.items():
7         # tokenize
8         desc = desc.split()
9         # convert to lower case
10        desc = [word.lower() for word in desc]
11        # remove punctuation from each token
12        desc = [w.translate(table) for w in desc]
13        # remove hanging 's' and 'a'
14        desc = [word for word in desc if len(word)>1]
15        # store as string
16        descriptions[key] = ' '.join(desc)
17
18 # clean descriptions
19 clean_descriptions(descriptions)
20 # summarize vocabulary
21 all_tokens = ' '.join(descriptions.values()).split()
22 vocabulary = set(all_tokens)
23 print('Vocabulary Size: %d' % len(vocabulary))

```

Finally, we save the dictionary of image identifiers and descriptions to a new file named *descriptions.txt*, with one image identifier and description per line.

Below defines the `save_doc()` function that given a dictionary containing the mapping of identifiers to descriptions and a filename, saves the mapping to file.

```

1 # save descriptions to file, one per line
2 def save_doc(descriptions, filename):
3     lines = list()
4     for key, desc in descriptions.items():
5         lines.append(key + ' ' + desc)
6     data = '\n'.join(lines)
7     file = open(filename, 'w')
8     file.write(data)
9     file.close()
10
11 # save descriptions
12 save_doc(descriptions, 'descriptions.txt')

```

Putting this all together, the complete listing is provided below.

```

1 import string
2
3 # load doc into memory
4 def load_doc(filename):
5     # open the file as read only
6     file = open(filename, 'r')
7     # read all text
8     text = file.read()
9     # close the file
10    file.close()
11    return text
12
13 # extract descriptions for images
14 def load_descriptions(doc):
15     mapping = dict()
16     # process lines
17     for line in doc.split('\n'):
18         # split line by white space
19         tokens = line.split()
20         if len(line) < 2:
21             continue

```

```

22     # take the first token as the image id, the rest as the description
23     image_id, image_desc = tokens[0], tokens[1:]
24     # remove filename from image id
25     image_id = image_id.split('.')[0]
26     # convert description tokens back to string
27     image_desc = ' '.join(image_desc)
28     # store the first description for each image
29     if image_id not in mapping:
30         mapping[image_id] = image_desc
31     return mapping
32
33 def clean_descriptions(descriptions):
34     # prepare translation table for removing punctuation
35     table = str.maketrans('', '', string.punctuation)
36     for key, desc in descriptions.items():
37         # tokenize
38         desc = desc.split()
39         # convert to lower case
40         desc = [word.lower() for word in desc]
41         # remove punctuation from each token
42         desc = [w.translate(table) for w in desc]
43         # remove hanging 's' and 'a'
44         desc = [word for word in desc if len(word)>1]
45         # store as string
46         descriptions[key] = ' '.join(desc)
47
48 # save descriptions to file, one per line
49 def save_doc(descriptions, filename):
50     lines = list()
51     for key, desc in descriptions.items():
52         lines.append(key + ' ' + desc)
53     data = '\n'.join(lines)
54     file = open(filename, 'w')
55     file.write(data)
56     file.close()
57
58 filename = 'Flickr8k_text/Flickr8k.token.txt'
59 # load descriptions
60 doc = load_doc(filename)
61 # parse descriptions
62 descriptions = load_descriptions(doc)
63 print('Loaded: %d ' % len(descriptions))
64 # clean descriptions
65 clean_descriptions(descriptions)
66 # summarize vocabulary
67 all_tokens = ' '.join(descriptions.values()).split()
68 vocabulary = set(all_tokens)
69 print('Vocabulary Size: %d' % len(vocabulary))
70 # save descriptions
71 save_doc(descriptions, 'descriptions.txt')

```

Running the example first prints the number of loaded photo descriptions (8,092) and the size of the clean vocabulary (4,484 words).

```

1 Loaded: 8092
2 Vocabulary Size: 4484

```

The clean descriptions are then written to `'descriptions.txt'`. Taking a look in the file, we can see that the descriptions are ready for modeling.

Taking a look in the file, we can see that the descriptions are ready for modeling.


```
1 3621647714_fc67ab2617 man is standing on snow with trees and mountains all around him
2 365128300_6966058139 group of people are rafting on river rapids
3 2751694538_fffa3d307d man and boy sit in the driver seat
4 537628742_146f2c24f8 little girl running in field
5 2320125735_27fe729948 black and brown dog with blue collar goes on alert by soccer ball in the g
6 ...
```

Preparing the Photos

We will use a pre-trained model to interpret the content of the photos.

There are many models to choose from. In this case, we will use the Oxford Visual Geometry Group or VGG model that won the ImageNet competition in 2014. Learn more about the model here:

- [Very Deep Convolutional Networks for Large-Scale Visual Recognition](#)

Keras provides this pre-trained model directly. Note, the first time you use this model, Keras will download the model weights from the Internet, which are about 500 Megabytes. This may take a few minutes depending on your internet connection.

We could use this model as part of a broader image caption model. The problem is, it is a large model and running each photo through the network every time we want to test a new language model configuration (downstream) is redundant.

Instead, we can pre-compute the “photo features” using the pre-trained model and save them to file. We can then load these features later and feed them into our model as the interpretation of a given photo in the dataset. It is no different to running the photo through the full VGG model, it is just that we will have done it once in advance.

This is an optimization that will make training our models faster and consume less memory.

We can load the VGG model in Keras using the VGG class. We will load the model without the top; this means without the layers at the end of the network that are used to interpret the features extracted from the input and turn them into a class prediction. We are not interested in the image net classification of the photos and we will train our own interpretation of the image features.

Keras also provides tools for reshaping the loaded photo into the preferred size for the model (e.g. 3 channel 224 x 224 pixel image).

Below is a function named *extract_features()* that given a directory name will load each photo, prepare it for VGG and collect the predicted features from the VGG model. The image features are a 3-dimensional array with the shape (7, 7, 512).

The function returns a dictionary of image identifier to image features.

```
1 # extract features from each photo in the directory
2 def extract_features(directory):
3     # load the model
4     in_layer = Input(shape=(224, 224, 3))
5     model = VGG16(include_top=False, input_tensor=in_layer)
```

```

6     print(model.summary())
7     # extract features from each photo
8     features = dict()
9     for name in listdir(directory):
10        # load an image from file
11        filename = directory + '/' + name
12        image = load_img(filename, target_size=(224, 224))
13        # convert the image pixels to a numpy array
14        image = img_to_array(image)
15        # reshape data for the model
16        image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))
17        # prepare the image for the VGG model
18        image = preprocess_input(image)
19        # get features
20        feature = model.predict(image, verbose=0)
21        # get image id
22        image_id = name.split('.')[0]
23        # store feature
24        features[image_id] = feature
25        print('>%s' % name)
26    return features

```

We can call this function to prepare the photo data for testing our models, then save the resulting dictionary to a file named 'features.pkl'.

The complete example is listed below.

```

1  from os import listdir
2  from pickle import dump
3  from keras.applications.vgg16 import VGG16
4  from keras.preprocessing.image import load_img
5  from keras.preprocessing.image import img_to_array
6  from keras.applications.vgg16 import preprocess_input
7  from keras.layers import Input
8
9  # extract features from each photo in the directory
10 def extract_features(directory):
11     # load the model
12     in_layer = Input(shape=(224, 224, 3))
13     model = VGG16(include_top=False, input_tensor=in_layer)
14     print(model.summary())
15     # extract features from each photo
16     features = dict()
17     for name in listdir(directory):
18        # load an image from file
19        filename = directory + '/' + name
20        image = load_img(filename, target_size=(224, 224))
21        # convert the image pixels to a numpy array
22        image = img_to_array(image)
23        # reshape data for the model
24        image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))
25        # prepare the image for the VGG model
26        image = preprocess_input(image)
27        # get features
28        feature = model.predict(image, verbose=0)
29        # get image id
30        image_id = name.split('.')[0]
31        # store feature
32        features[image_id] = feature
33        print('>%s' % name)
34    return features
35
36 # extract features from all images

```



```
37 directory = 'Flicker8k_Dataset'
38 features = extract_features(directory)
39 print('Extracted Features: %d' % len(features))
40 # save to file
41 dump(features, open('features.pkl', 'wb'))
```

Running this data preparation step may take a while depending on your hardware, perhaps one hour on the CPU with a modern workstation.

At the end of the run, you will have the extracted features stored in *'features.pkl'* for later use.

Baseline Caption Generation Model

In this section, we will define a baseline model for generating captions for photos and how to evaluate it so that it can be compared to variations on this baseline.

This section is divided into 5 parts:

1. Load Data.
2. Fit Model.
3. Evaluate Model.
4. Complete Example
5. “A” versus “A” Test
6. Generate Photo Captions

1. Load Data

We are not going to fit the model on all of the caption data, or even on a large sample of the data.

In this tutorial, we are interested in quickly testing a suite of different configurations of a caption model to see what works on this data. That means we need the evaluation of one model configuration to happen quickly. Toward this end, we will train the models on 100 photographs and captions, then evaluate them on both the training dataset and on a new test set of 100 photographs and captions.

First, we need to load a pre-defined subset of photographs. The provided dataset has separate sets for train, test, and development, which are really just different groups of photo identifiers. We will load the development set and use the first 100 identifiers for train and the second 100 (e.g. from 100 to 200) as the test set.

The function *load_set()* below will load a pre-defined set of identifiers, and we will call it with the *'Flickr_8k.devImages.txt'* filename as an argument.

```
1 # load a pre-defined list of photo identifiers
2 def load_set(filename):
3     doc = load_doc(filename)
4     dataset = list()
5     # process line by line
6     for line in doc.split('\n'):
7         # skip empty lines
8         if len(line) < 1:
9             continue
```

```

10         # get the image identifier
11         identifier = line.split('.')[0]
12         dataset.append(identifier)
13     return set(dataset)

```

Next, we need to split the set into train and test sets.

We will start by ordering the identifiers by sorting them to ensure we always split them consistently across machines and runs, then take the first 100 for train and the next 100 for test.

The `train_test_split()` function below will create this split given the loaded set of identifiers as input.

```

1 # split a dataset into train/test elements
2 def train_test_split(dataset):
3     # order keys so the split is consistent
4     ordered = sorted(dataset)
5     # return split dataset as two new sets
6     return set(ordered[:100]), set(ordered[100:200])

```

Now, we can load the photo descriptions using the pre-defined set of train or test identifiers.

Below is the function `load_clean_descriptions()` that loads the cleaned text descriptions from 'descriptions.txt' for a given set of identifiers and returns a dictionary of identifier to text.

The model we will develop will generate a caption given a photo, and the caption will be generated one word at a time. The sequence of previously generated words will be provided as input. Therefore, we will need a "first word" to kick-off the generation process and a 'last word' to signal the end of the caption. We will use the strings 'startseq' and 'endseq' for this purpose.

```

1 # load clean descriptions into memory
2 def load_clean_descriptions(filename, dataset):
3     # load document
4     doc = load_doc(filename)
5     descriptions = dict()
6     for line in doc.split('\n'):
7         # split line by white space
8         tokens = line.split()
9         # split id from description
10        image_id, image_desc = tokens[0], tokens[1:]
11        # skip images not in the set
12        if image_id in dataset:
13            # store
14            descriptions[image_id] = 'startseq ' + ' '.join(image_desc) + ' endseq'
15    return descriptions

```

Next, we can load the photo features for a given dataset.

Below defines a function named `load_photo_features()` that loads the entire set of photo descriptions, then returns the subset of interest for a given set of photo identifiers. This is not very efficient as the loaded dictionary of all photo features is about 700 Megabytes. Nevertheless, this will get us up and running quickly.

Note, if you have a better approach, share it in the comments below.

```

1 # load photo features

```

```

2 def load_photo_features(filename, dataset):
3     # load all features
4     all_features = load(open(filename, 'rb'))
5     # filter features
6     features = {k: all_features[k] for k in dataset}
7     return features

```

We can pause here and test everything developed so far.

The complete code example is listed below.

```

1  from pickle import load
2
3  # load doc into memory
4  def load_doc(filename):
5      # open the file as read only
6      file = open(filename, 'r')
7      # read all text
8      text = file.read()
9      # close the file
10     file.close()
11     return text
12
13 # load a pre-defined list of photo identifiers
14 def load_set(filename):
15     doc = load_doc(filename)
16     dataset = list()
17     # process line by line
18     for line in doc.split('\n'):
19         # skip empty lines
20         if len(line) < 1:
21             continue
22         # get the image identifier
23         identifier = line.split('.')[0]
24         dataset.append(identifier)
25     return set(dataset)
26
27 # split a dataset into train/test elements
28 def train_test_split(dataset):
29     # order keys so the split is consistent
30     ordered = sorted(dataset)
31     # return split dataset as two new sets
32     return set(ordered[:100]), set(ordered[100:200])
33
34 # load clean descriptions into memory
35 def load_clean_descriptions(filename, dataset):
36     # load document
37     doc = load_doc(filename)
38     descriptions = dict()
39     for line in doc.split('\n'):
40         # split line by white space
41         tokens = line.split()
42         # split id from description
43         image_id, image_desc = tokens[0], tokens[1:]
44         # skip images not in the set
45         if image_id in dataset:
46             # store
47             descriptions[image_id] = 'startseq ' + ' '.join(image_desc) + ' endseq'
48     return descriptions
49
50 # load photo features
51 def load_photo_features(filename, dataset):
52     # load all features
53     all_features = load(open(filename, 'rb'))

```

```

54     # filter features
55     features = {k: all_features[k] for k in dataset}
56     return features
57
58 # load dev set
59 filename = 'Flickr8k_text/Flickr_8k.devImages.txt'
60 dataset = load_set(filename)
61 print('Dataset: %d' % len(dataset))
62 # train-test split
63 train, test = train_test_split(dataset)
64 print('Train=%d, Test=%d' % (len(train), len(test)))
65 # descriptions
66 train_descriptions = load_clean_descriptions('descriptions.txt', train)
67 test_descriptions = load_clean_descriptions('descriptions.txt', test)
68 print('Descriptions: train=%d, test=%d' % (len(train_descriptions), len(test_descriptions)))
69 # photo features
70 train_features = load_photo_features('features.pkl', train)
71 test_features = load_photo_features('features.pkl', test)
72 print('Photos: train=%d, test=%d' % (len(train_features), len(test_features)))

```

Running this example first loads the 1,000 photo identifiers in the development dataset. A train and test set is selected and used to filter the set of clean photo descriptions and prepared image features.

We are nearly there.

```

1 Dataset: 1,000
2 Train=100, Test=100
3 Descriptions: train=100, test=100
4 Photos: train=100, test=100

```

The description text will need to be encoded to numbers before it can be presented to the model as in input or compared to the model's predictions.

The first step in encoding the data is to create a consistent mapping from words to unique integer values. Keras provides the `Tokenizer` class that can learn this mapping from the loaded description data.

Below defines the `create_tokenizer()` that will fit a `Tokenizer` given the loaded photo description text.

```

1 # fit a tokenizer given caption descriptions
2 def create_tokenizer(descriptions):
3     lines = list(descriptions.values())
4     tokenizer = Tokenizer()
5     tokenizer.fit_on_texts(lines)
6     return tokenizer
7
8 # prepare tokenizer
9 tokenizer = create_tokenizer(descriptions)
10 vocab_size = len(tokenizer.word_index) + 1
11 print('Vocabulary Size: %d' % vocab_size)

```

We can now encode the text.

Each description will be split into words. The model will be provided one word and the photo and generate the next word. Then the first two words of the description will be provided to the model as input with the image to generate the next word. This is how the model will be trained.

For example, the input sequence “*little girl running in field*” would be split into 6 input-output pairs to train the model:

1	X1,	X2 (text sequence),	y (word)
2	photo	startseq,	little
3	photo	startseq, little,	girl
4	photo	startseq, little, girl,	running
5	photo	startseq, little, girl, running,	in
6	photo	startseq, little, girl, running, in,	field
7	photo	startseq, little, girl, running, in, field, endseq	

Later when the model is used to generate descriptions, the generated words will be concatenated and recursively provided as input to generate a caption for an image.

The function below named `create_sequences()` given the tokenizer, a single clean description, the features for a photo, and the maximum description length will prepare a set of input-output pairs for training a model. Calling this function will return `X1` and `X2` for the arrays of image data and input sequence data and the `y` value for the output word.

The input sequences are integer encoded and the output word is one-hot encoded to represent the probability distribution of the expected word across the whole vocabulary of possible words.

```

1  # create sequences of images, input sequences and output words for an image
2  def create_sequences(tokenizer, desc, image, max_length):
3      Ximages, XSeq, y = list(), list(), list()
4      vocab_size = len(tokenizer.word_index) + 1
5      # integer encode the description
6      seq = tokenizer.texts_to_sequences([desc])[0]
7      # split one sequence into multiple X,y pairs
8      for i in range(1, len(seq)):
9          # select
10         in_seq, out_seq = seq[:i], seq[i]
11         # pad input sequence
12         in_seq = pad_sequences([in_seq], maxlen=max_length)[0]
13         # encode output sequence
14         out_seq = to_categorical([out_seq], num_classes=vocab_size)[0]
15         # store
16         Ximages.append(image)
17         XSeq.append(in_seq)
18         y.append(out_seq)
19     # Ximages, XSeq, y = array(Ximages), array(XSeq), array(y)
20     return [Ximages, XSeq, y]
```

2. Fit Model

We are nearly ready to fit the model.

Parts of the model have already been discussed, but let's re-iterate.

The model is based on the example laid out in the paper “[Show and Tell: A Neural Image Caption Generator](#)”, 2015.

The model involves three parts:

- **Photo Feature Extractor.** This is a 16-layer VGG model pre-trained on the ImageNet dataset. We have pre-processed the photos with a the VGG model (without the top) and will use the extracted features predicted by this model as input.
- **Sequence Processor.** This is a word embedding layer for handling the text input, followed by an LSTM layer. The LSTM output is interpreted by a Dense layer one output at a time.
- **Interpreter (for lack of a better name).** Both the feature extractor and sequence processor output a fixed-length vector that is the length of a maximum sequence. These are concatenated together and processed by an LSTM and Dense layer before a final prediction is made.

A conservative number of neurons is used in the base model. Specifically, a 128 Dense layer after the feature extractor, a 50-dimensionality word embedding followed by a 256 unit LSTM and 128 neuron Dense after the sequence processor, and finally a 500 unit LSTM followed by a 500 neuron Dense at the end of the network.

The model predicts a probability distribution across the vocabulary, therefore a softmax activation function is used and a categorical cross entropy loss function is minimized while fitting the network.

The function `define_model()` defines the baseline model, given the size of the vocabulary and the maximum length of photo descriptions. The Keras functional API is used to define the model as it provides the flexibility needed to define a model that takes two input streams and combines them.

```

1 # define the captioning model
2 def define_model(vocab_size, max_length):
3     # feature extractor (encoder)
4     inputs1 = Input(shape=(7, 7, 512))
5     fe1 = GlobalMaxPooling2D()(inputs1)
6     fe2 = Dense(128, activation='relu')(fe1)
7     fe3 = RepeatVector(max_length)(fe2)
8     # embedding
9     inputs2 = Input(shape=(max_length,))
10    emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
11    emb3 = LSTM(256, return_sequences=True)(emb2)
12    emb4 = TimeDistributed(Dense(128, activation='relu'))(emb3)
13    # merge inputs
14    merged = concatenate([fe3, emb4])
15    # language model (decoder)
16    lm2 = LSTM(500)(merged)
17    lm3 = Dense(500, activation='relu')(lm2)
18    outputs = Dense(vocab_size, activation='softmax')(lm3)
19    # tie it together [image, seq] [word]
20    model = Model(inputs=[inputs1, inputs2], outputs=outputs)
21    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
22    print(model.summary())
23    plot_model(model, show_shapes=True, to_file='plot.png')
24    return model

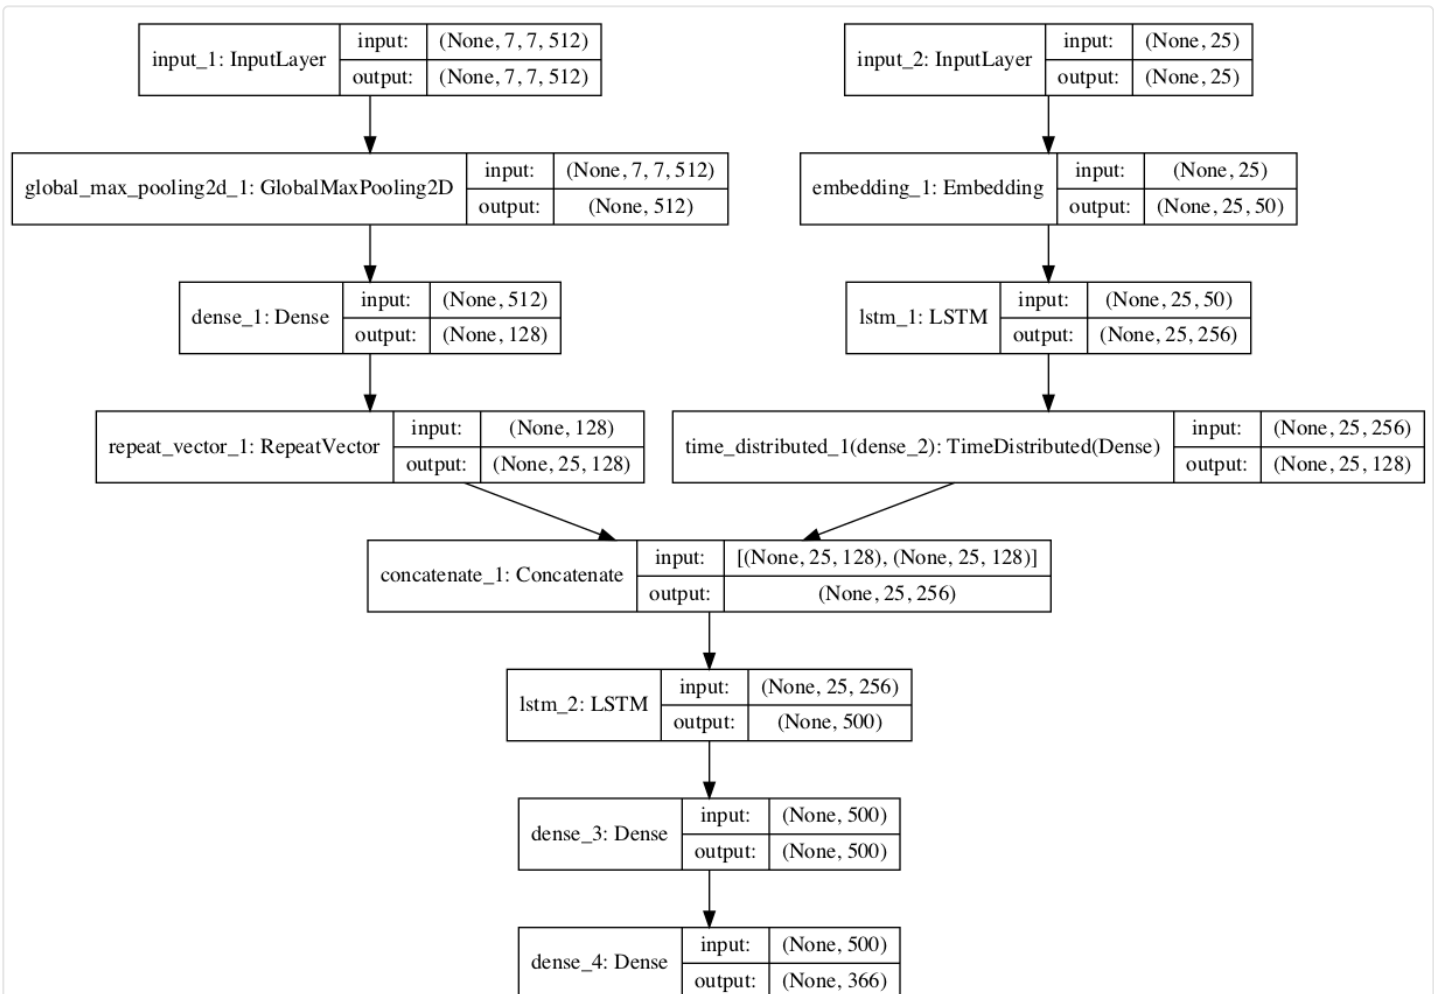
```

To get a sense for the structure of the model, specifically the shapes of the layers, see the summary listed below.

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 7, 7, 512)	0	
input_2 (InputLayer)	(None, 25)	0	

7				
8	global_max_pooling2d_1 (GlobalMa	(None, 512)	0	input_1[0][0]
9				
10	embedding_1 (Embedding)	(None, 25, 50)	18300	input_2[0][0]
11				
12	dense_1 (Dense)	(None, 128)	65664	global_max_pooling2d_1[0][0]
13				
14	lstm_1 (LSTM)	(None, 25, 256)	314368	embedding_1[0][0]
15				
16	repeat_vector_1 (RepeatVector)	(None, 25, 128)	0	dense_1[0][0]
17				
18	time_distributed_1 (TimeDistribu	(None, 25, 128)	32896	lstm_1[0][0]
19				
20	concatenate_1 (Concatenate)	(None, 25, 256)	0	repeat_vector_1[0][0] time_distributed_1[0][0]
21				
22				
23	lstm_2 (LSTM)	(None, 500)	1514000	concatenate_1[0][0]
24				
25	dense_3 (Dense)	(None, 500)	250500	lstm_2[0][0]
26				
27	dense_4 (Dense)	(None, 366)	183366	dense_3[0][0]
28				
29	Total params: 2,379,094			
30	Trainable params: 2,379,094			
31	Non-trainable params: 0			
32				

We also create a plot to visualize the structure of the network that better helps understand the two streams of input.



We will train the model using a data generator. This is strictly not required given that the captions and extracted photo features can probably fit into memory as a single dataset. Nevertheless, it is good practice for when you come to train the final model on the entire dataset.

A generator will yield a result when called. In Keras, it will yield a single batch of input-output samples that are used to estimate the error gradient and update the model weights.

The function `data_generator()` defines the data generator, given a dictionary of loaded photo descriptions, photo features, the tokenizer for integer encoding sequences, and the maximum sequence length in the dataset.

The generator loops forever and keeps yielding batches of input-output pairs when asked. We also have a `n_step` parameter that allows us to tune how many images worth of input-output pairs to generate for each batch. The average sequence has 10 words, that is 10 input-output pairs, and a good batch size might be 30 samples, which is about 2-to-3 images worth.

```

1 # data generator, intended to be used in a call to model.fit_generator()
2 def data_generator(descriptions, features, tokenizer, max_length, n_step):
3     # loop until we finish training
4     while 1:
5         # loop over photo identifiers in the dataset
6         keys = list(descriptions.keys())
7         for i in range(0, len(keys), n_step):
8             Ximages, XSeq, y = list(), list(), list()
9             for j in range(i, min(len(keys), i+n_step)):
10                 image_id = keys[j]
11                 # retrieve photo feature input
12                 image = features[image_id][0]
13                 # retrieve text input
14                 desc = descriptions[image_id]
15                 # generate input-output pairs
16                 in_img, in_seq, out_word = create_sequences(tokenizer, desc, image, max_length)
17                 for k in range(len(in_img)):
18                     Ximages.append(in_img[k])
19                     XSeq.append(in_seq[k])
20                     y.append(out_word[k])
21                 # yield this batch of samples to the model
22                 yield [[array(Ximages), array(XSeq)], array(y)]

```

The model can be fit by calling `fit_generator()` and passing it to the data generator, along with all of the parameters needed. When fitting the model, we can also specify the number of batches to run per epoch and the number of epochs.

```

1 model.fit_generator(data_generator(train_descriptions, train_features, tokenizer, max_length, n_

```

For these experiments, we will use 2 images per batch, 50 batches (or 100 images) per epoch, and 50 training epochs. You can experiment with different configurations in your own experiments.

3. Evaluate Model

Now that we know how to prepare the data and define a model, we must define a test harness to evaluate a given model.

We will evaluate a model by training it on the dataset, generating descriptions for all photos in the training dataset, evaluating those predictions with a cost function, and then repeating this evaluation process multiple times.

The outcome will be a distribution of skill scores for the model that we can summarize by calculating the mean and standard deviation. This is the preferred way to evaluate deep learning models. See this post:

- [How to Evaluate the Skill of Deep Learning Models](#)

First, we need to be able to generate a description for a photo using a trained model.

This involves passing in the start description token `'startseq'`, generating one word, then calling the model recursively with generated words as input until the end of sequence token is reached `'endseq'` or the maximum description length is reached.

The function below named `generate_desc()` implements this behavior and generates a textual description given a trained model, and a given prepared photo as input. It calls the function `word_for_id()` in order to map an integer prediction back to a word.

```

1  # map an integer to a word
2  def word_for_id(integer, tokenizer):
3      for word, index in tokenizer.word_index.items():
4          if index == integer:
5              return word
6      return None
7
8  # generate a description for an image
9  def generate_desc(model, tokenizer, photo, max_length):
10     # seed the generation process
11     in_text = 'startseq'
12     # iterate over the whole length of the sequence
13     for i in range(max_length):
14         # integer encode input sequence
15         sequence = tokenizer.texts_to_sequences([in_text])[0]
16         # pad input
17         sequence = pad_sequences([sequence], maxlen=max_length)
18         # predict next word
19         yhat = model.predict([photo, sequence], verbose=0)
20         # convert probability to integer
21         yhat = argmax(yhat)
22         # map integer to word
23         word = word_for_id(yhat, tokenizer)
24         # stop if we cannot map the word
25         if word is None:
26             break
27         # append as input for generating the next word
28         in_text += ' ' + word
29         # stop if we predict the end of the sequence
30         if word == 'endseq':
31             break
32     return in_text

```

We will generate predictions for all photos in the training dataset and in the test dataset.

The function below named `evaluate_model()` will evaluate a trained model against a given dataset of photo descriptions and photo features. The actual and predicted descriptions are collected and evaluated

collectively using the corpus BLEU score that summarizes how close the generated text is to the expected text.

```

1 # evaluate the skill of the model
2 def evaluate_model(model, descriptions, photos, tokenizer, max_length):
3     actual, predicted = list(), list()
4     # step over the whole set
5     for key, desc in descriptions.items():
6         # generate description
7         yhat = generate_desc(model, tokenizer, photos[key], max_length)
8         # store actual and predicted
9         actual.append([desc.split()])
10        predicted.append(yhat.split())
11    # calculate BLEU score
12    bleu = corpus_bleu(actual, predicted)
13    return bleu

```

BLEU scores are used in text translation for evaluating translated text against one or more reference translations. We do in fact have access to multiple reference descriptions for each image that we could compare to, but for simplicity, we will use the first description for each photo in the dataset (e.g. the cleaned version).

You can learn more about the BLEU score here:

- [BLEU \(bilingual evaluation understudy\) on Wikipedia](#)

The NLTK Python library implements the BLEU score calculation in the `corpus_bleu()` function. A higher score close to 1.0 is better, a score closer to zero is worse.

Finally, all we need to do is define, fit, and evaluate the model multiple times in a loop then report the final average score.

Ideally, we would repeat the experiment 30 times or more, but this will take too long for our small test harness. Instead, will evaluate the model 3 times. It will be faster, but the mean score will have higher variance.

Below defines the model evaluation loop. At the end of the run, the distribution of BLEU scores for the train and test sets are saved to a file.

```

1 # run experiment
2 train_results, test_results = list(), list()
3 for i in range(n_repeats):
4     # define the model
5     model = define_model(vocab_size, max_length)
6     # fit model
7     model.fit_generator(data_generator(train_descriptions, train_features, tokenizer, max_length),
8     # evaluate model on training data
9     train_score = evaluate_model(model, train_descriptions, train_features, tokenizer, max_length)
10    test_score = evaluate_model(model, test_descriptions, test_features, tokenizer, max_length)
11    # store
12    train_results.append(train_score)
13    test_results.append(test_score)
14    print('>%d: train=%f test=%f' % ((i+1), train_score, test_score))
15 # save results to file
16 df = DataFrame()
17 df['train'] = train_results

```

```

18 df['test'] = test_results
19 print(df.describe())
20 df.to_csv(model_name+'.csv', index=False)

```

We parameterize the run as follows, allowing us to name each run and save the result to separate files.

```

1 # define experiment
2 model_name = 'baseline1'
3 verbose = 2
4 n_epochs = 50
5 n_photos_per_update = 2
6 n_batches_per_epoch = int(len(train) / n_photos_per_update)
7 n_repeats = 3

```

4. Complete Example

The complete example is listed below.

```

1 from os import listdir
2 from numpy import array
3 from numpy import argmax
4 from pandas import DataFrame
5 from nltk.translate.bleu_score import corpus_bleu
6 from pickle import load
7
8 from keras.preprocessing.text import Tokenizer
9 from keras.preprocessing.sequence import pad_sequences
10 from keras.utils import to_categorical
11 from keras.preprocessing.image import load_img
12 from keras.preprocessing.image import img_to_array
13 from keras.applications.vgg16 import preprocess_input
14 from keras.applications.vgg16 import VGG16
15 from keras.utils import plot_model
16 from keras.models import Model
17 from keras.layers import Input
18 from keras.layers import Dense
19 from keras.layers import Flatten
20 from keras.layers import LSTM
21 from keras.layers import RepeatVector
22 from keras.layers import TimeDistributed
23 from keras.layers import Embedding
24 from keras.layers.merge import concatenate
25 from keras.layers.pooling import GlobalMaxPooling2D
26
27 # load doc into memory
28 def load_doc(filename):
29     # open the file as read only
30     file = open(filename, 'r')
31     # read all text
32     text = file.read()
33     # close the file
34     file.close()
35     return text
36
37 # load a pre-defined list of photo identifiers
38 def load_set(filename):
39     doc = load_doc(filename)
40     dataset = list()
41     # process line by line
42     for line in doc.split('\n'):
43         # skip empty lines
44         if len(line) < 1:
45             continue

```

```

46         # get the image identifier
47         identifier = line.split('.')[0]
48         dataset.append(identifier)
49     return set(dataset)
50
51 # split a dataset into train/test elements
52 def train_test_split(dataset):
53     # order keys so the split is consistent
54     ordered = sorted(dataset)
55     # return split dataset as two new sets
56     return set(ordered[:100]), set(ordered[100:200])
57
58 # load clean descriptions into memory
59 def load_clean_descriptions(filename, dataset):
60     # load document
61     doc = load_doc(filename)
62     descriptions = dict()
63     for line in doc.split('\n'):
64         # split line by white space
65         tokens = line.split()
66         # split id from description
67         image_id, image_desc = tokens[0], tokens[1:]
68         # skip images not in the set
69         if image_id in dataset:
70             # store
71             descriptions[image_id] = 'startseq ' + ' '.join(image_desc) + ' endseq'
72     return descriptions
73
74 # load photo features
75 def load_photo_features(filename, dataset):
76     # load all features
77     all_features = load(open(filename, 'rb'))
78     # filter features
79     features = {k: all_features[k] for k in dataset}
80     return features
81
82 # fit a tokenizer given caption descriptions
83 def create_tokenizer(descriptions):
84     lines = list(descriptions.values())
85     tokenizer = Tokenizer()
86     tokenizer.fit_on_texts(lines)
87     return tokenizer
88
89 # create sequences of images, input sequences and output words for an image
90 def create_sequences(tokenizer, desc, image, max_length):
91     Ximages, XSeq, y = list(), list(), list()
92     vocab_size = len(tokenizer.word_index) + 1
93     # integer encode the description
94     seq = tokenizer.texts_to_sequences([desc])[0]
95     # split one sequence into multiple X,y pairs
96     for i in range(1, len(seq)):
97         # select
98         in_seq, out_seq = seq[:i], seq[i]
99         # pad input sequence
100        in_seq = pad_sequences([in_seq], maxlen=max_length)[0]
101        # encode output sequence
102        out_seq = to_categorical([out_seq], num_classes=vocab_size)[0]
103        # store
104        Ximages.append(image)
105        XSeq.append(in_seq)
106        y.append(out_seq)
107    # Ximages, XSeq, y = array(Ximages), array(XSeq), array(y)
108    return [Ximages, XSeq, y]
109
110 # define the captioning model

```



```

111 def define_model(vocab_size, max_length):
112     # feature extractor (encoder)
113     inputs1 = Input(shape=(7, 7, 512))
114     fe1 = GlobalMaxPooling2D()(inputs1)
115     fe2 = Dense(128, activation='relu')(fe1)
116     fe3 = RepeatVector(max_length)(fe2)
117     # embedding
118     inputs2 = Input(shape=(max_length,))
119     emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
120     emb3 = LSTM(256, return_sequences=True)(emb2)
121     emb4 = TimeDistributed(Dense(128, activation='relu'))(emb3)
122     # merge inputs
123     merged = concatenate([fe3, emb4])
124     # language model (decoder)
125     lm2 = LSTM(500)(merged)
126     lm3 = Dense(500, activation='relu')(lm2)
127     outputs = Dense(vocab_size, activation='softmax')(lm3)
128     # tie it together [image, seq] [word]
129     model = Model(inputs=[inputs1, inputs2], outputs=outputs)
130     model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
131     print(model.summary())
132     plot_model(model, show_shapes=True, to_file='plot.png')
133     return model
134
135 # data generator, intended to be used in a call to model.fit_generator()
136 def data_generator(descriptions, features, tokenizer, max_length, n_step):
137     # loop until we finish training
138     while 1:
139         # loop over photo identifiers in the dataset
140         keys = list(descriptions.keys())
141         for i in range(0, len(keys), n_step):
142             Ximages, XSeq, y = list(), list(), list()
143             for j in range(i, min(len(keys), i+n_step)):
144                 image_id = keys[j]
145                 # retrieve photo feature input
146                 image = features[image_id][0]
147                 # retrieve text input
148                 desc = descriptions[image_id]
149                 # generate input-output pairs
150                 in_img, in_seq, out_word = create_sequences(tokenizer, desc, image, max_length)
151                 for k in range(len(in_img)):
152                     Ximages.append(in_img[k])
153                     XSeq.append(in_seq[k])
154                     y.append(out_word[k])
155             # yield this batch of samples to the model
156             yield [[array(Ximages), array(XSeq)], array(y)]
157
158 # map an integer to a word
159 def word_for_id(integer, tokenizer):
160     for word, index in tokenizer.word_index.items():
161         if index == integer:
162             return word
163     return None
164
165 # generate a description for an image
166 def generate_desc(model, tokenizer, photo, max_length):
167     # seed the generation process
168     in_text = 'startseq'
169     # iterate over the whole length of the sequence
170     for i in range(max_length):
171         # integer encode input sequence
172         sequence = tokenizer.texts_to_sequences([in_text])[0]
173         # pad input
174         sequence = pad_sequences([sequence], maxlen=max_length)
175         # predict next word

```

```

176     yhat = model.predict([photo,sequence], verbose=0)
177     # convert probability to integer
178     yhat = argmax(yhat)
179     # map integer to word
180     word = word_for_id(yhat, tokenizer)
181     # stop if we cannot map the word
182     if word is None:
183         break
184     # append as input for generating the next word
185     in_text += ' ' + word
186     # stop if we predict the end of the sequence
187     if word == 'endseq':
188         break
189     return in_text
190
191 # evaluate the skill of the model
192 def evaluate_model(model, descriptions, photos, tokenizer, max_length):
193     actual, predicted = list(), list()
194     # step over the whole set
195     for key, desc in descriptions.items():
196         # generate description
197         yhat = generate_desc(model, tokenizer, photos[key], max_length)
198         # store actual and predicted
199         actual.append([desc.split()])
200         predicted.append(yhat.split())
201     # calculate BLEU score
202     bleu = corpus_bleu(actual, predicted)
203     return bleu
204
205 # load dev set
206 filename = 'Flickr8k_text/Flickr_8k.devImages.txt'
207 dataset = load_set(filename)
208 print('Dataset: %d' % len(dataset))
209 # train-test split
210 train, test = train_test_split(dataset)
211 # descriptions
212 train_descriptions = load_clean_descriptions('descriptions.txt', train)
213 test_descriptions = load_clean_descriptions('descriptions.txt', test)
214 print('Descriptions: train=%d, test=%d' % (len(train_descriptions), len(test_descriptions)))
215 # photo features
216 train_features = load_photo_features('features.pkl', train)
217 test_features = load_photo_features('features.pkl', test)
218 print('Photos: train=%d, test=%d' % (len(train_features), len(test_features)))
219 # prepare tokenizer
220 tokenizer = create_tokenizer(train_descriptions)
221 vocab_size = len(tokenizer.word_index) + 1
222 print('Vocabulary Size: %d' % vocab_size)
223 # determine the maximum sequence length
224 max_length = max(len(s.split()) for s in list(train_descriptions.values()))
225 print('Description Length: %d' % max_length)
226
227 # define experiment
228 model_name = 'baseline1'
229 verbose = 2
230 n_epochs = 50
231 n_photos_per_update = 2
232 n_batches_per_epoch = int(len(train) / n_photos_per_update)
233 n_repeats = 3
234
235 # run experiment
236 train_results, test_results = list(), list()
237 for i in range(n_repeats):
238     # define the model
239     model = define_model(vocab_size, max_length)
240     # fit model

```

```

241 model.fit_generator(data_generator(train_descriptions, train_features, tokenizer, max_length,
242 # evaluate model on training data
243 train_score = evaluate_model(model, train_descriptions, train_features, tokenizer, max_length,
244 test_score = evaluate_model(model, test_descriptions, test_features, tokenizer, max_length,
245 # store
246 train_results.append(train_score)
247 test_results.append(test_score)
248 print('>%d: train=%f test=%f' % ((i+1), train_score, test_score))
249 # save results to file
250 df = DataFrame()
251 df['train'] = train_results
252 df['test'] = test_results
253 print(df.describe())
254 df.to_csv(model_name+'.csv', index=False)

```

Running the example first prints summary statistics for the loaded training data.

```

1 Dataset: 1,000
2 Descriptions: train=100, test=100
3 Photos: train=100, test=100
4 Vocabulary Size: 366
5 Description Length: 25

```

The example should take about 20 minutes on GPU hardware, a little longer on CPU hardware.

At the end of the run, a mean BLEU of 0.06 is reported on the training set and 0.04 on the test set. Results are stored in *baseline1.csv*.

1		train	test
2	count	3.000000	3.000000
3	mean	0.060617	0.040978
4	std	0.023498	0.025105
5	min	0.042882	0.012101
6	25%	0.047291	0.032658
7	50%	0.051701	0.053215
8	75%	0.069484	0.055416
9	max	0.087268	0.057617

This provides a baseline model for comparison to alternate configurations.

“A” versus “A” Test

Before we start testing variations of the model, it is important to get an idea of whether or not the test harness is stable.

That is, whether the summarizing skill of the model over 5 runs is sufficient to control for the stochastic nature of the model.

We can get an idea of this by running the experiment again in what is called an A vs A test in A/B testing land. We would expect to get an equivalent result if we ran the same experiment again; if we don't, perhaps additional repeats would be required to control for the stochastic nature of the method and on the dataset.

Below are the results from a second run of the algorithm.

1		train	test
2	count	3.000000	3.000000
3	mean	0.036902	0.043003

4	std	0.020281	0.017295
5	min	0.018522	0.026055
6	25%	0.026023	0.034192
7	50%	0.033525	0.042329
8	75%	0.046093	0.051477
9	max	0.058660	0.060624

We can see that the run gets a very similar mean and standard deviation BLEU scores. Specifically, a mean BLEU of 0.03 vs 0.06 on train and 0.04 to 0.04 for test.

The harness is a little noisy, but stable enough for comparison.

Is the model any good?

Generate Photo Captions

We expect the model is under-trained and maybe even under provisioned, but can it generate any kind of readable text at all?

It is important that the baseline model have some modicum of capability so that we can relate the BLEU scores of the baseline to an idea of what kind of quality of descriptions are being generated.

Let's train a single model and generate a few descriptions from the train and test sets as a sanity check.

Change the number of repeats to 1 and the name of the run to `'baseline_generate'`.

```
1 model_name = 'baseline_generate'
2 n_repeats = 1
```

Then update the `evaluate_model()` function to only evaluate the first 5 photos in the dataset and print the descriptions, as follows.

```
1 # evaluate the skill of the model
2 def evaluate_model(model, descriptions, photos, tokenizer, max_length):
3     actual, predicted = list(), list()
4     # step over the whole set
5     for key, desc in descriptions.items():
6         # generate description
7         yhat = generate_desc(model, tokenizer, photos[key], max_length)
8         # store actual and predicted
9         actual.append([desc.split()])
10        predicted.append(yhat.split())
11        print('Actual:    %s' % desc)
12        print('Predicted: %s' % yhat)
13        if len(actual) >= 5:
14            break
15    # calculate BLEU score
16    bleu = corpus_bleu(actual, predicted)
17    return bleu
```

Re-run the example.

You should see results for the train set like the following (the specific results will vary given the stochastic nature of the algorithm):

```
1 Actual:    startseq boy bites hard into treat while he sits outside endseq
2 Predicted: startseq boy boy while while he while outside endseq
3
4 Actual:    startseq man in field backed by american flags endseq
5 Predicted: startseq man in in standing city endseq
6
7 Actual:    startseq two girls are walking down dirt road in park endseq
8 Predicted: startseq man walking down down road in endseq
9
10 Actual:   startseq girl laying on the tree with boy kneeling before her endseq
11 Predicted: startseq boy while in up up up up water endseq
12
13 Actual:   startseq boy in striped shirt is jumping in front of water fountain endseq
14 Predicted: startseq man is is shirt is on on on on on bike endseq
```

You should see results on the test dataset as follows:

```
1 Actual:    startseq three people are looking into photographic equipment endseq
2 Predicted: startseq boy racer on on on on on bike endseq
3
4 Actual:    startseq boy is leaning on chair whilst another boy pulls him around with rope endseq
5 Predicted: startseq girl in playing on on on sword endseq
6
7 Actual:    startseq black and brown dog jumping in midair near field endseq
8 Predicted: startseq dog dog running running running and dog in grass endseq
9
10 Actual:   startseq dog places his head on man face endseq
11 Predicted: startseq brown dog to to to to to to to ball endseq
12
13 Actual:   startseq man in green hat is someplace up high endseq
14 Predicted: startseq man in up up his waves endseq
```

We can see that the descriptions are not perfect, some are a little rough, but generally the model is generating somewhat readable text. A good starting point for improvement.

Next, let's look at some experiments to vary the size or capacity of different sub-models.

Network Size Parameters

In this section, we will see how gross variations to the network structure impact model skill.

We will look at the following aspects of the model size:

1. Size of the fixed-vector output from the 'encoders'.
2. Size of the sequence encoder model.
3. Size of the language model.

Let's dive in.

Size of Fixed-Length Vector

In the baseline model, the photo feature extractor and the text sequence encoder both output a 128 element vector. These vectors are then concatenated to be processed by the language model.

The 128 element vector from each sub-model contains everything known about the input sequence and photo. We can vary the size of this vector to see if it impacts model skill

First, we can decrease the size by half from 128 elements to 64 elements.

```

1 # define the captioning model
2 def define_model(vocab_size, max_length):
3     # feature extractor (encoder)
4     inputs1 = Input(shape=(7, 7, 512))
5     fe1 = GlobalMaxPooling2D()(inputs1)
6     fe2 = Dense(64, activation='relu')(fe1)
7     fe3 = RepeatVector(max_length)(fe2)
8     # embedding
9     inputs2 = Input(shape=(max_length,))
10    emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
11    emb3 = LSTM(256, return_sequences=True)(emb2)
12    emb4 = TimeDistributed(Dense(64, activation='relu'))(emb3)
13    # merge inputs
14    merged = concatenate([fe3, emb4])
15    # language model (decoder)
16    lm2 = LSTM(500)(merged)
17    lm3 = Dense(500, activation='relu')(lm2)
18    outputs = Dense(vocab_size, activation='softmax')(lm3)
19    # tie it together [image, seq] [word]
20    model = Model(inputs=[inputs1, inputs2], outputs=outputs)
21    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
22    return model

```

We will name this model 'size_sm_fixed_vec'.

```
1 model_name = 'size_sm_fixed_vec'
```

Running this experiment produces the following BLEU scores, perhaps a small gain over baseline on the test set.

1		train	test
2	count	3.000000	3.000000
3	mean	0.204421	0.063148
4	std	0.026992	0.003264
5	min	0.174769	0.059391
6	25%	0.192849	0.062074
7	50%	0.210929	0.064757
8	75%	0.219246	0.065026
9	max	0.227564	0.065295

We can also double the size of the fixed-length vector from 128 to 256 units.

```

1 # define the captioning model
2 def define_model(vocab_size, max_length):
3     # feature extractor (encoder)
4     inputs1 = Input(shape=(7, 7, 512))
5     fe1 = GlobalMaxPooling2D()(inputs1)
6     fe2 = Dense(256, activation='relu')(fe1)
7     fe3 = RepeatVector(max_length)(fe2)
8     # embedding
9     inputs2 = Input(shape=(max_length,))
10    emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
11    emb3 = LSTM(256, return_sequences=True)(emb2)
12    emb4 = TimeDistributed(Dense(256, activation='relu'))(emb3)
13    # merge inputs
14    merged = concatenate([fe3, emb4])

```



```

15 # language model (decoder)
16 lm2 = LSTM(500)(merged)
17 lm3 = Dense(500, activation='relu')(lm2)
18 outputs = Dense(vocab_size, activation='softmax')(lm3)
19 # tie it together [image, seq] [word]
20 model = Model(inputs=[inputs1, inputs2], outputs=outputs)
21 model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
22 return model

```

We will name this configuration 'size_lg_fixed_vec'.

```
1 model_name = 'size_lg_fixed_vec'
```

Running this experiment shows BLEU scores suggesting that the model is not better off.

It is possible that with more data and/or longer training, we may see a different story.

	train	test
count	3.000000	3.000000
mean	0.023517	0.027813
std	0.009951	0.010525
min	0.012037	0.021737
25%	0.020435	0.021737
50%	0.028833	0.021737
75%	0.029257	0.030852
max	0.029682	0.039966

Sequence Encoder Size

We can call the sub-model that interprets the input sequence of words generated so far as the sequence encoder.

First, we can try to see if decreasing the representational capacity of the sequence encoder impacts model skill. We can reduce the number of memory units in the LSTM layer from 256 to 128.

```

1 # define the captioning model
2 def define_model(vocab_size, max_length):
3     # feature extractor (encoder)
4     inputs1 = Input(shape=(7, 7, 512))
5     fe1 = GlobalMaxPooling2D()(inputs1)
6     fe2 = Dense(128, activation='relu')(fe1)
7     fe3 = RepeatVector(max_length)(fe2)
8     # embedding
9     inputs2 = Input(shape=(max_length,))
10    emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
11    emb3 = LSTM(128, return_sequences=True)(emb2)
12    emb4 = TimeDistributed(Dense(128, activation='relu'))(emb3)
13    # merge inputs
14    merged = concatenate([fe3, emb4])
15    # language model (decoder)
16    lm2 = LSTM(500)(merged)
17    lm3 = Dense(500, activation='relu')(lm2)
18    outputs = Dense(vocab_size, activation='softmax')(lm3)
19    # tie it together [image, seq] [word]
20    model = Model(inputs=[inputs1, inputs2], outputs=outputs)
21    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
22    return model
23
24 model_name = 'size_sm_seq_model'

```

Running this example, we can see perhaps a small bump on both train and test over baseline. This might be an artifact of the small training set size.

	train	test
1 count	3.000000	3.000000
2 mean	0.074944	0.053917
3 std	0.014263	0.013264
4 min	0.066292	0.039142
5 25%	0.066713	0.048476
6 50%	0.067134	0.057810
7 75%	0.079270	0.061304
8 max	0.091406	0.064799

Going the other way, we can double the number of LSTM layers from one to two and see if that makes a dramatic difference.

```

1 # define the captioning model
2 def define_model(vocab_size, max_length):
3     # feature extractor (encoder)
4     inputs1 = Input(shape=(7, 7, 512))
5     fe1 = GlobalMaxPooling2D()(inputs1)
6     fe2 = Dense(128, activation='relu')(fe1)
7     fe3 = RepeatVector(max_length)(fe2)
8     # embedding
9     inputs2 = Input(shape=(max_length,))
10    emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
11    emb3 = LSTM(256, return_sequences=True)(emb2)
12    emb4 = LSTM(256, return_sequences=True)(emb3)
13    emb5 = TimeDistributed(Dense(128, activation='relu'))(emb4)
14    # merge inputs
15    merged = concatenate([fe3, emb5])
16    # language model (decoder)
17    lm2 = LSTM(500)(merged)
18    lm3 = Dense(500, activation='relu')(lm2)
19    outputs = Dense(vocab_size, activation='softmax')(lm3)
20    # tie it together [image, seq] [word]
21    model = Model(inputs=[inputs1, inputs2], outputs=outputs)
22    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
23    return model
24
25 model_name = 'size_lg_seq_model'

```

Running this experiment shows a decent bump in BLEU on both train and test sets.

	train	test
1 count	3.000000	3.000000
2 mean	0.094937	0.096970
3 std	0.022394	0.079270
4 min	0.069151	0.046722
5 25%	0.087656	0.051279
6 50%	0.106161	0.055836
7 75%	0.107830	0.122094
8 max	0.109499	0.188351

We can also try to increase the representational capacity of the word embedding by doubling it from 50-dimensions to 100-dimensions.

```

1 # define the captioning model
2 def define_model(vocab_size, max_length):
3     # feature extractor (encoder)
4     inputs1 = Input(shape=(7, 7, 512))

```

```

5 fe1 = GlobalMaxPooling2D()(inputs1)
6 fe2 = Dense(128, activation='relu')(fe1)
7 fe3 = RepeatVector(max_length)(fe2)
8 # embedding
9 inputs2 = Input(shape=(max_length,))
10 emb2 = Embedding(vocab_size, 100, mask_zero=True)(inputs2)
11 emb3 = LSTM(256, return_sequences=True)(emb2)
12 emb4 = TimeDistributed(Dense(128, activation='relu'))(emb3)
13 # merge inputs
14 merged = concatenate([fe3, emb4])
15 # language model (decoder)
16 lm2 = LSTM(500)(merged)
17 lm3 = Dense(500, activation='relu')(lm2)
18 outputs = Dense(vocab_size, activation='softmax')(lm3)
19 # tie it together [image, seq] [word]
20 model = Model(inputs=[inputs1, inputs2], outputs=outputs)
21 model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
22 return model
23
24 model_name = 'size_em_seq_model'

```

We see a large movement on the training dataset, but perhaps little movement on the test dataset.

1	count	3.000000	3.000000
2	mean	0.112743	0.050935
3	std	0.017136	0.006860
4	min	0.096121	0.043741
5	25%	0.103940	0.047701
6	50%	0.111759	0.051661
7	75%	0.121055	0.054533
8	max	0.130350	0.057404

Size of Language Model

We can refer to the model that learns from the concatenated sequence and photo feature input as the language model. It is responsible for generating words.

First, we can look at the impact on model skill by cutting the LSTM and dense layers from 500 to 256 neurons.

```

1 # define the captioning model
2 def define_model(vocab_size, max_length):
3     # feature extractor (encoder)
4     inputs1 = Input(shape=(7, 7, 512))
5     fe1 = GlobalMaxPooling2D()(inputs1)
6     fe2 = Dense(128, activation='relu')(fe1)
7     fe3 = RepeatVector(max_length)(fe2)
8     # embedding
9     inputs2 = Input(shape=(max_length,))
10    emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
11    emb3 = LSTM(256, return_sequences=True)(emb2)
12    emb4 = TimeDistributed(Dense(128, activation='relu'))(emb3)
13    # merge inputs
14    merged = concatenate([fe3, emb4])
15    # language model (decoder)
16    lm2 = LSTM(256)(merged)
17    lm3 = Dense(256, activation='relu')(lm2)
18    outputs = Dense(vocab_size, activation='softmax')(lm3)
19    # tie it together [image, seq] [word]
20    model = Model(inputs=[inputs1, inputs2], outputs=outputs)
21    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

```

```

22     return model
23
24 model_name = 'size_sm_lang_model'

```

We can see that this has a small positive effect on BLEU for both training and test datasets, again, likely related to the small size of the datasets.

	train	test
count	3.000000	3.000000
mean	0.063632	0.056059
std	0.018521	0.009064
min	0.045127	0.048916
25%	0.054363	0.050961
50%	0.063599	0.053005
75%	0.072884	0.059630
max	0.082169	0.066256

We can also look at the impact of doubling the capacity of the language model by adding a second LSTM layer of the same size.

```

1  # define the captioning model
2  def define_model(vocab_size, max_length):
3      # feature extractor (encoder)
4      inputs1 = Input(shape=(7, 7, 512))
5      fe1 = GlobalMaxPooling2D()(inputs1)
6      fe2 = Dense(128, activation='relu')(fe1)
7      fe3 = RepeatVector(max_length)(fe2)
8      # embedding
9      inputs2 = Input(shape=(max_length,))
10     emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
11     emb3 = LSTM(256, return_sequences=True)(emb2)
12     emb4 = TimeDistributed(Dense(128, activation='relu'))(emb3)
13     # merge inputs
14     merged = concatenate([fe3, emb4])
15     # language model (decoder)
16     lm2 = LSTM(500, return_sequences=True)(merged)
17     lm3 = LSTM(500)(lm2)
18     lm4 = Dense(500, activation='relu')(lm3)
19     outputs = Dense(vocab_size, activation='softmax')(lm4)
20     # tie it together [image, seq] [word]
21     model = Model(inputs=[inputs1, inputs2], outputs=outputs)
22     model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
23     return model
24
25 model_name = 'size_lg_lang_model'

```

Again, we see minor movements in BLEU, perhaps an artifact of noise and dataset size. The improvement on the test

The improvement on the test dataset may be a good sign. This might be a change worth exploring.

	train	test
count	3.000000	3.000000
mean	0.043838	0.067658
std	0.037580	0.045813
min	0.017990	0.015757
25%	0.022284	0.050252
50%	0.026578	0.084748
75%	0.056763	0.093608
max	0.086948	0.102469

Tuning model size on a much smaller dataset is challenging.

Configuring the Feature Extraction Model

The use of the pre-trained VGG16 model provides some additional points of configuration.

The baseline model removed the top from the VGG model, including a global max pooling layer, which then feeds into an encoding of the features to a 128 element vector.

In this section, we will look at the following modifications to the baseline model:

1. Using a global average pooling layer after the VGG model.
2. Not using any global pooling.

Global Average Pooling

We can replace the GlobalMaxPooling2D layer with a GlobalAveragePooling2D to achieve average pooling.

Global average pooling was developed to reduce overfitting for image classification problems, but may offer some benefit in interpreting the features extracted from the image.

For more on global average pooling, see the paper:

- [Network In Network](#), 2013.

The updated `define_model()` function and experiment name are listed below.

```

1  # define the captioning model
2  def define_model(vocab_size, max_length):
3      # feature extractor (encoder)
4      inputs1 = Input(shape=(7, 7, 512))
5      fe1 = GlobalAveragePooling2D()(inputs1)
6      fe2 = Dense(128, activation='relu')(fe1)
7      fe3 = RepeatVector(max_length)(fe2)
8      # embedding
9      inputs2 = Input(shape=(max_length,))
10     emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
11     emb3 = LSTM(256, return_sequences=True)(emb2)
12     emb4 = TimeDistributed(Dense(128, activation='relu'))(emb3)
13     # merge inputs
14     merged = concatenate([fe3, emb4])
15     # language model (decoder)
16     lm2 = LSTM(500)(merged)
17     lm3 = Dense(500, activation='relu')(lm2)
18     outputs = Dense(vocab_size, activation='softmax')(lm3)
19     # tie it together [image, seq] [word]
20     model = Model(inputs=[inputs1, inputs2], outputs=outputs)
21     model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
22     return model
23
24 model_name = 'fe_avg_pool'

```

The results suggest a dramatic improvement on the training dataset, which may be a sign of overfitting. We also see a small lift on test skill. This might be a change worth exploring.

We also see a small lift on test skill. This might be a change worth exploring.

	train	test
1 count	3.000000	3.000000
2 mean	0.834627	0.060847
3 std	0.083259	0.040463
4 min	0.745074	0.017705
5 25%	0.797096	0.042294
6 50%	0.849118	0.066884
7 75%	0.879404	0.082418
8 max	0.909690	0.097952

No Pooling

We can remove the GlobalMaxPooling2D and flatten the 3D photo feature and feed it directly into a Dense layer.

I would not expect this to be a good model design, but it is worth testing this assumption.

```

1 # define the captioning model
2 def define_model(vocab_size, max_length):
3     # feature extractor (encoder)
4     inputs1 = Input(shape=(7, 7, 512))
5     fe1 = Flatten()(inputs1)
6     fe2 = Dense(128, activation='relu')(fe1)
7     fe3 = RepeatVector(max_length)(fe2)
8     # embedding
9     inputs2 = Input(shape=(max_length,))
10    emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
11    emb3 = LSTM(256, return_sequences=True)(emb2)
12    emb4 = TimeDistributed(Dense(128, activation='relu'))(emb3)
13    # merge inputs
14    merged = concatenate([fe3, emb4])
15    # language model (decoder)
16    lm2 = LSTM(500)(merged)
17    lm3 = Dense(500, activation='relu')(lm2)
18    outputs = Dense(vocab_size, activation='softmax')(lm3)
19    # tie it together [image, seq] [word]
20    model = Model(inputs=[inputs1, inputs2], outputs=outputs)
21    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
22    return model
23
24 model_name = 'fe_flat'

```

Surprisingly, we see a small lift on training data and a large lift on test data. This is surprising (to me) and may be worth further investigation.

	train	test
1 count	3.000000	3.000000
2 mean	0.055988	0.135231
3 std	0.017566	0.079714
4 min	0.038605	0.044177
5 25%	0.047116	0.106633
6 50%	0.055627	0.169089
7 75%	0.064679	0.180758
8 max	0.073731	0.192428

We can try repeating this experiment and provide more capacity for interpreting the extracted photo features. A new Dense layer with 500 neurons is added after the Flatten layer.

```

1 # define the captioning model
2 def define_model(vocab_size, max_length):
3     # feature extractor (encoder)
4     inputs1 = Input(shape=(7, 7, 512))
5     fe1 = Flatten()(inputs1)
6     fe2 = Dense(500, activation='relu')(fe1)
7     fe3 = Dense(128, activation='relu')(fe2)
8     fe4 = RepeatVector(max_length)(fe3)
9     # embedding
10    inputs2 = Input(shape=(max_length,))
11    emb2 = Embedding(vocab_size, 50, mask_zero=True)(inputs2)
12    emb3 = LSTM(256, return_sequences=True)(emb2)
13    emb4 = TimeDistributed(Dense(128, activation='relu'))(emb3)
14    # merge inputs
15    merged = concatenate([fe4, emb4])
16    # language model (decoder)
17    lm2 = LSTM(500)(merged)
18    lm3 = Dense(500, activation='relu')(lm2)
19    outputs = Dense(vocab_size, activation='softmax')(lm3)
20    # tie it together [image, seq] [word]
21    model = Model(inputs=[inputs1, inputs2], outputs=outputs)
22    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
23    return model
24
25 model_name = 'fe_flat2'

```

This results in a less impressive change and perhaps worse BLEU results on the test dataset.

1		train	test
2	count	3.000000	3.000000
3	mean	0.060126	0.029487
4	std	0.030300	0.013205
5	min	0.031235	0.020850
6	25%	0.044359	0.021887
7	50%	0.057483	0.022923
8	75%	0.074572	0.033805
9	max	0.091661	0.044688

Word Embedding Models

A key part of the model is the sequence learning model that must interpret the sequence of words generated so far for a photo.

At the input to this sub-model is a word embedding and a good way to improve a word embedding over learning it from scratch as part of the model (as in the baseline model) is to use pre-trained word embeddings.

In this section, we will explore the impact of using a pre-trained word embedding on the model. Specifically:

1. Training a Word2Vec Model
2. Training a Word2Vec Model + Fine Tuning

Trained word2vec Embedding

An efficient learning algorithm for pre-training a word embedding from a corpus of text is the word2vec algorithm.

You can learn more about the word2vec algorithm here:

- [Word2Vec Google Code Project](#)

We can use this algorithm to train a new standalone set of word vectors using the cleaned photo descriptions in the dataset.

The [Gensim library](#) provides access to an implementation of the algorithm that we can use to pre-train the embedding.

First, we must load the clean photo descriptions for the training dataset, as before.

Next, we can fit the word2vec model on all of the clean descriptions. We should note that this includes more descriptions than the 50 used in the training dataset. A fairer model for these experiments should only be trained on those descriptions in the training dataset.

Once fit, we can save the words and word vectors to an ASCII file, perhaps for later inspection or visualization.

```
1 # train word2vec model
2 lines = [s.split() for s in train_descriptions.values()]
3 model = Word2Vec(lines, size=100, window=5, workers=8, min_count=1)
4 # summarize vocabulary size in model
5 words = list(model.wv.vocab)
6 print('Vocabulary size: %d' % len(words))
7
8 # save model in ASCII (word2vec) format
9 filename = 'custom_embedding.txt'
10 model.wv.save_word2vec_format(filename, binary=False)
```

The word embedding is saved to the file `'custom_embedding.txt'`.

Now, we can load the embedding into memory, retrieve only the word vectors for the words in our vocabulary, then save them to a new file.

```
1 # load the whole embedding into memory
2 embedding = dict()
3 file = open('custom_embedding.txt')
4 for line in file:
5     values = line.split()
6     word = values[0]
7     coefs = asarray(values[1:], dtype='float32')
8     embedding[word] = coefs
9 file.close()
10 print('Embedding Size: %d' % len(embedding))
11
12 # summarize vocabulary
13 all_tokens = ' '.join(train_descriptions.values()).split()
14 vocabulary = set(all_tokens)
15 print('Vocabulary Size: %d' % len(vocabulary))
16
17 # get the vectors for words in our vocab
```

```

18 cust_embedding = dict()
19 for word in vocabulary:
20     # check if word in embedding
21     if word not in embedding:
22         continue
23     cust_embedding[word] = embedding[word]
24 print('Custom Embedding %d' % len(cust_embedding))
25
26 # save
27 dump(cust_embedding, open('word2vec_embedding.pkl', 'wb'))
28 print('Saved Embedding')

```

The complete example is listed below.

```

1  # prepare word vectors for captioning model
2
3  from numpy import asarray
4  from pickle import dump
5  from gensim.models import Word2Vec
6
7  # load doc into memory
8  def load_doc(filename):
9      # open the file as read only
10     file = open(filename, 'r')
11     # read all text
12     text = file.read()
13     # close the file
14     file.close()
15     return text
16
17 # load a pre-defined list of photo identifiers
18 def load_set(filename):
19     doc = load_doc(filename)
20     dataset = list()
21     # process line by line
22     for line in doc.split('\n'):
23         # skip empty lines
24         if len(line) < 1:
25             continue
26         # get the image identifier
27         identifier = line.split('.')[0]
28         dataset.append(identifier)
29     return set(dataset)
30
31 # split a dataset into train/test elements
32 def train_test_split(dataset):
33     # order keys so the split is consistent
34     ordered = sorted(dataset)
35     # return split dataset as two new sets
36     return set(ordered[:100]), set(ordered[100:200])
37
38 # load clean descriptions into memory
39 def load_clean_descriptions(filename, dataset):
40     # load document
41     doc = load_doc(filename)
42     descriptions = dict()
43     for line in doc.split('\n'):
44         # split line by white space
45         tokens = line.split()
46         # split id from description
47         image_id, image_desc = tokens[0], tokens[1:]
48         # skip images not in the set
49         if image_id in dataset:
50             # store

```

```

51         descriptions[image_id] = 'startseq ' + ' '.join(image_desc) + ' endseq'
52     return descriptions
53
54 # load dev set
55 filename = 'Flickr8k_text/Flickr_8k.devImages.txt'
56 dataset = load_set(filename)
57 print('Dataset: %d' % len(dataset))
58 # train-test split
59 train, test = train_test_split(dataset)
60 print('Train=%d, Test=%d' % (len(train), len(test)))
61 # descriptions
62 train_descriptions = load_clean_descriptions('descriptions.txt', train)
63 print('Descriptions: train=%d' % len(train_descriptions))
64
65 # train word2vec model
66 lines = [s.split() for s in train_descriptions.values()]
67 model = Word2Vec(lines, size=100, window=5, workers=8, min_count=1)
68 # summarize vocabulary size in model
69 words = list(model.wv.vocab)
70 print('Vocabulary size: %d' % len(words))
71
72 # save model in ASCII (word2vec) format
73 filename = 'custom_embedding.txt'
74 model.wv.save_word2vec_format(filename, binary=False)
75
76 # load the whole embedding into memory
77 embedding = dict()
78 file = open('custom_embedding.txt')
79 for line in file:
80     values = line.split()
81     word = values[0]
82     coefs = asarray(values[1:], dtype='float32')
83     embedding[word] = coefs
84 file.close()
85 print('Embedding Size: %d' % len(embedding))
86
87 # summarize vocabulary
88 all_tokens = ' '.join(train_descriptions.values()).split()
89 vocabulary = set(all_tokens)
90 print('Vocabulary Size: %d' % len(vocabulary))
91
92 # get the vectors for words in our vocab
93 cust_embedding = dict()
94 for word in vocabulary:
95     # check if word in embedding
96     if word not in embedding:
97         continue
98     cust_embedding[word] = embedding[word]
99 print('Custom Embedding %d' % len(cust_embedding))
100
101 # save
102 dump(cust_embedding, open('word2vec_embedding.pkl', 'wb'))
103 print('Saved Embedding')

```

Running this example creates a new dictionary mapping of word-to-word vectors stored in the file 'word2vec_embedding.pkl'.

```

1 Dataset: 1000
2 Train=100, Test=100
3 Descriptions: train=100
4 Vocabulary size: 365
5 Embedding Size: 366
6 Vocabulary Size: 365
7 Custom Embedding 365

```

8 Saved Embedding

Next, we can load this embedding and use the word vectors as the fixed weights in an Embedding layer.

Below provides the `load_embedding()` function that loads the custom word2vec embedding and returns the new Embedding layer for use in the model.

```

1 # load a word embedding
2 def load_embedding(tokenizer, vocab_size, max_length):
3     # load the tokenizer
4     embedding = load(open('word2vec_embedding.pkl', 'rb'))
5     dimensions = 100
6     trainable = False
7     # create a weight matrix for words in training docs
8     weights = zeros((vocab_size, dimensions))
9     # walk words in order of tokenizer vocab to ensure vectors are in the right index
10    for word, i in tokenizer.word_index.items():
11        if word not in embedding:
12            continue
13        weights[i] = embedding[word]
14    layer = Embedding(vocab_size, dimensions, weights=[weights], input_length=max_length, trainable=trainable)
15    return layer

```

We can use it in our model by calling the function directly from our `define_model()` function.

```

1 # define the captioning model
2 def define_model(tokenizer, vocab_size, max_length):
3     # feature extractor (encoder)
4     inputs1 = Input(shape=(7, 7, 512))
5     fe1 = GlobalMaxPooling2D()(inputs1)
6     fe2 = Dense(128, activation='relu')(fe1)
7     fe3 = RepeatVector(max_length)(fe2)
8     # embedding
9     inputs2 = Input(shape=(max_length,))
10    emb2 = load_embedding(tokenizer, vocab_size, max_length)(inputs2)
11    emb3 = LSTM(256, return_sequences=True)(emb2)
12    emb4 = TimeDistributed(Dense(128, activation='relu'))(emb3)
13    # merge inputs
14    merged = concatenate([fe3, emb4])
15    # language model (decoder)
16    lm2 = LSTM(500)(merged)
17    lm3 = Dense(500, activation='relu')(lm2)
18    outputs = Dense(vocab_size, activation='softmax')(lm3)
19    # tie it together [image, seq] [word]
20    model = Model(inputs=[inputs1, inputs2], outputs=outputs)
21    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
22    return model
23
24 model_name = 'seq_w2v_fixed'

```

We can see some lift on the training dataset, perhaps no real notable change on the test dataset.

1		train	test
2	count	3.000000	3.000000
3	mean	0.096780	0.047540
4	std	0.055073	0.008445
5	min	0.033511	0.038340
6	25%	0.078186	0.043840
7	50%	0.122861	0.049341
8	75%	0.128414	0.052140
9	max	0.133967	0.054939

Trained word2vec Embedding with Fine Tuning

We can repeat the previous experiment and allow the model to tune the word vectors while fitting the model.

The updated `load_embedding()` function that permits the embedding layer to be fine-tuned is listed below.

```

1  # load a word embedding
2  def load_embedding(tokenizer, vocab_size, max_length):
3      # load the tokenizer
4      embedding = load(open('word2vec_embedding.pkl', 'rb'))
5      dimensions = 100
6      trainable = True
7      # create a weight matrix for words in training docs
8      weights = zeros((vocab_size, dimensions))
9      # walk words in order of tokenizer vocab to ensure vectors are in the right index
10     for word, i in tokenizer.word_index.items():
11         if word not in embedding:
12             continue
13         weights[i] = embedding[word]
14     layer = Embedding(vocab_size, dimensions, weights=[weights], input_length=max_length, trainable=trainable)
15     return layer
16
17 model_name = 'seq_w2v_tuned'

```

Again, we do not see much difference in using these pre-trained word embedding vectors over the baseline model.

1		train	test
2	count	3.000000	3.000000
3	mean	0.065297	0.042712
4	std	0.080194	0.007697
5	min	0.017675	0.034593
6	25%	0.019003	0.039117
7	50%	0.020332	0.043641
8	75%	0.089108	0.046772
9	max	0.157885	0.049904

Analysis of Results

We have performed a few experiments on a very small sample (1.6%) from the Flickr8k training dataset of 8,000 photos.

It is possible that the sample is too small, that the models were not trained for long enough, and that 3 repeats of each model results in too much variance. These aspects can also be tested by evaluating by designing experiments such as:


1. Does model skill scale with the size of the dataset?
2. Do more epochs result in better skill?
3. Do more repeats result in a skill with less variance?

Nevertheless, we have some ideas on how we might configure a model for the fuller dataset.

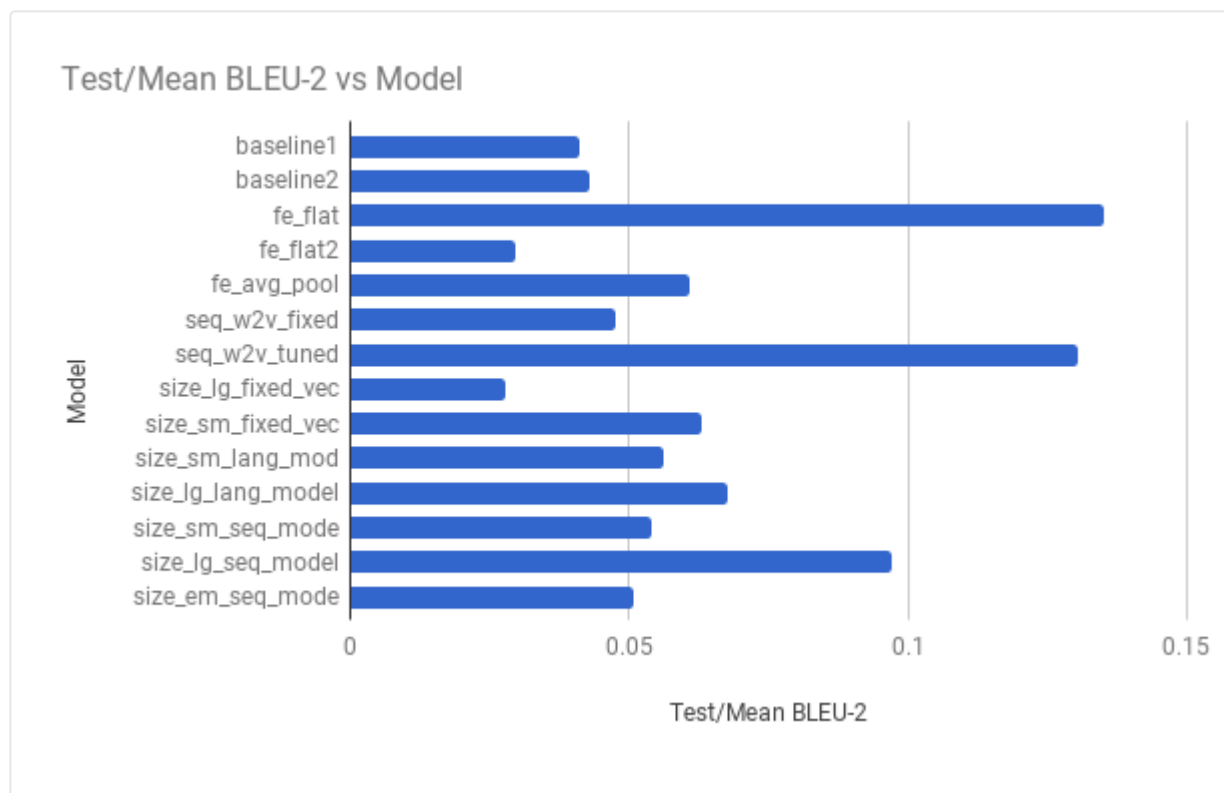
Below is a summary of the mean results from the experiments performed in this tutorial.

It is helpful to review a graph of the results. If we had more repeats, a box and whisker plot for each distribution of scores might be a good visualization. Here we use a simple bar graph. Remember, that larger BLEU scores are better.

Results on the training dataset:



Results on the test dataset:



From just looking at the mean results on the test dataset, we can suggest:

- Perhaps pooling is not required after the photo feature extractor (fe_flat at 0.135231).
- Perhaps average pooling offers an advantage over max pooling after the photo feature extractor (fe_avg_pool at 0.060847).
- Perhaps a smaller sized fixed-length vector after the sub-models is a good idea (size_sm_fixed_vec at 0.063148).
- Perhaps adding more layers to the language model offers some benefit (size_lg_lang_model at 0.067658).
- Perhaps adding more layers to the sequence model offers some benefit (size_lg_seq_model at 0.09697).


I would also recommend exploring combinations of these suggestions.

We can also review the distribution of results.

Below is some code to load the saved results from each experiment and create a box-and-whisker plot of results on the train and test sets for review.


```
1 from os import listdir
2 from pandas import read_csv
3 from pandas import DataFrame
4 from matplotlib import pyplot
5
6 # load all .csv results into a dataframe
7 train, test = DataFrame(), DataFrame()
8 directory = 'results'
9 for name in listdir(directory):
10     if not name.endswith('.csv'):
11         continue
12     filename = directory + '/' + name
13     data = read_csv(filename, header=0)
14     experiment = name.split('.')[0]
15     train[experiment] = data['train']
16     test[experiment] = data['test']
17
18 # plot results on train
19 train.boxplot(vert=False)
20 pyplot.show()
21 # plot results on test
22 test.boxplot(vert=False)
23 pyplot.show()
```

Distribution of results on the training dataset.



Box and Whisker Plot of Experiment vs Model Skill on the Training Dataset

Distribution of results on the test dataset.



Box and Whisker Plot of Experiment vs Model Skill on the Test Dataset

A review of these distributions suggests:

- The spread on the flat results is large; perhaps going with average pooling might be safer.
- The spread on the larger language model is large and skewed in the wrong/risky direction.
- The spread on the larger sequence model is large and skewed in the right direction.
- There may be some benefit in a smaller fixed-length vector size.

I would expect increasing repeats to 5, 10, or 30 would tighten up these distributions somewhat.

Further Reading

This section provides more resources on the topic if you are looking go deeper.

Papers

- [Show and Tell: A Neural Image Caption Generator](#), 2015.
- [Show, Attend and Tell: Neural Image Caption Generation with Visual Attention](#), 2016.
- [Network In Network](#), 2013.

Related Captioning Projects

- [caption_generator](#): An image captioning project
- [Keras Image Caption](#)
- [Neural Image Captioning \(NIC\)](#)
- [Keras deep learning for image caption retrieval](#)
- [DataLab Cup 2: Image Captioning](#)

Other

- [Framing image description as a ranking task: data, models and evaluation metrics](#).

API

- [Keras VGG16 API](#)
- [Gensim word2vec API](#)

Summary

In this tutorial, you discovered how you can use a small sample of the photo captioning dataset to explore different model designs.

Specifically, you learned:

- How to prepare data for photo captioning modeling.
- How to design a baseline and test harness to evaluate the skill of models and control for their stochastic nature.
- How to evaluate properties like model skill, feature extraction model, and word embeddings in order to lift model skill.


What experiments can you think up?

What else have you tried?

What are the best results you can get on the train and test dataset?

Let me know in the comments below.

Develop Deep Learning models for Text Data Today!



Develop Your Own Text models in Minutes

...with just a few lines of python code

Discover how in my new Ebook:

Deep Learning for Natural Language Processing

It provides **self-study tutorials** on topics like:

Bag-of-Words, Word Embedding, Language Models, Caption Generation, Text Translation and much more...

Finally Bring Deep Learning to your Natural Language Processing Projects

Skip the Academics. Just Results.

[Click to learn more.](#)



About Jason Brownlee

Jason Brownlee, PhD is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee →](#)

< A Gentle Introduction to Deep Learning Caption Generation Models

How to Develop a Deep Learning Photo Caption Generator from Scratch >

41 Responses to *How to Use Small Experiments to Develop a Caption Generation Model in Keras*

Emil November 24, 2017 at 6:39 am #

REPLY ↩

Hats off, another ace tutorial!

I'm curious how the TimeDistributed layer impacts the data before the concatenation. Is it possible to skip it? Also, is there a reason you are using VGG instead of the InceptionResNetV2 class other than memory/compute constraints.

Thanks!



Jason Brownlee November 24, 2017 at 9:52 am #

REPLY ↩

I chose VGG because it is smaller and simpler. You can use anything you wish.

You can skip the TimeDistributed as Dense can support time steps now I believe. I like it in there as it reminds me what is going on (e.g. outputting time steps).



Alex November 24, 2017 at 6:16 pm #

REPLY ↩

Ho jason, why don t you reset the LSTM states between the inputs related to different images? Ad they are not related to the same sequences.



Jason Brownlee November 25, 2017 at 10:14 am #

REPLY ↩

For speed of training.

It's a great suggestion though, try it and see if it lifts skill! Let me know how you go.



Alex November 28, 2017 at 1:39 am #

REPLY ↩

Thanks! In order to try this way, should I set stateful=True (avoiding the LSTM to reset itself automatically) and manually run `model.reset_states()` before training a single batch? (each batch is related to the sequence of a single image).



Jason Brownlee November 28, 2017 at 8:38 am #

REPLY ↩

Yes.

Emil December 13, 2017 at 12:10 am #

REPLY ↩

What's the logic of the +1 when you are creating the vocab len: "vocab_size = len(tokenizer.word_index) + 1"? Is it to leave remove for the 0?

Thanks



Jason Brownlee December 13, 2017 at 5:39 am #

REPLY ↩

Good question, to make space for 0 – words in the vocab start at 1.



Yang Cheng March 12, 2018 at 8:50 pm #

REPLY ↩

Does it mean we leave index 0 for 'endseq' token?



Jason Brownlee March 13, 2018 at 6:27 am #

REPLY ↩

No, the start and end tokens are legit parts of the problem. The model must specify when the sequence has ended.



xiaolian December 23, 2017 at 1:54 am #

REPLY ↩

i got a error : Error when checking input: expected input_11 to have 4 dimensions, but got array with shape (28, 4096)



Jason Brownlee December 23, 2017 at 5:21 am #

REPLY ↩

Are you able to confirm that your libraries are up to date and that you copied all of the code from the post?



xiaolian December 23, 2017 at 1:40 pm #

REPLY ↩

can i get the code in github?



Jason Brownlee December 24, 2017 at 4:51 am #

REPLY ↩

The code is on the post, why do you need it on github?



Niels February 24, 2018 at 12:41 am #

REPLY ↩

Hello Jason.

Thank you very much for your article.

I have a practical question regarding training an encoder decoder network.

So basically I have images of serial numbers and I want to predict the full serial number. (E.G 018F6176)
So if I train a network to predict the next character I would practically have to create a for loop predicting until I reach max length or the stop word.

This I get. However how can i structure my data?

What I have is the following

data is a numpy float array of shape (nb_samples width, height, nb_channels=3,).

labels is a numpy integer array of shape (nb_samples, max_caption_len)

So if I were to construct a similar dataset with the structure you get from `create_sequences()` my number of samples would increase and how do I ensure that images are loaded so as the sequence for each images comes in the correct order(or does that matter?)

Hopefully this makes sense to you.

Best regards

Niels



Jason Brownlee February 24, 2018 at 9:17 am #

REPLY ↩

Great question.

The caption model is the approach you want to use I believe.

See this post, specifically the section titled “ord-By-Word Model”:

<https://machinelearningmastery.com/prepare-photo-caption-dataset-training-deep-learning-model/>

It will show you exactly how to prepare your data and how to think about it.



Niels February 26, 2018 at 8:01 pm #

REPLY ↩

Thank you. This is exactly what I was looking for. One question though. Does the order of input matter, so let's say all samples for image 1 comes in order, all samples for image 2 next etc. Because then I would have to create a batch generator, however I would like to avoid that. My initial thought is it doesn't really matter since you are just training the algorithm to recognize what comes next regarding previous input. Is that a correct assumption?

Best Regards

Niels



Jason Brownlee February 27, 2018 at 6:25 am #

REPLY ↩

I think so, I think all samples for a give photo should be together by intuition (LSTMs have memory across the batch), but testing all assumptions is a good idea.

Niels February 27, 2018 at 9:27 pm #



Okay. I created a model both with the data_generator and randomly splitting in train,test(image order scrambled). The tran, test model did not seem to converge(max 17% validation accuracy). However the data_generator model(where order is preserved) reaches 99.5% validation accuracy , so that's pretty awesome.

Thank you for your time 😊



Jason Brownlee February 28, 2018 at 6:03 am #

Very nice. Thanks for running the experiment!



Binay March 7, 2018 at 12:18 pm #

REPLY ↩

Why and How the dimension of input_1 is (7, 7, 512)



Jason Brownlee March 7, 2018 at 3:06 pm #

REPLY ↩

That is the shape that we saved the extracted features from the photos.



Steven March 23, 2018 at 9:03 pm #

REPLY ↩

Thanks for your tutorials, Jason. I can run this tutorial in contrary with the tutorial on the same dataset <https://machinelearningmastery.com/develop-a-deep-learning-caption-generation-model-in-python/> where memory errors occur.

I compared both examples and found that the feature extraction of the images is different. This example the features are a 3-dimensional array with the shape (7, 7, 512) the other example the result is a 1-dimensional 4,096 element vector. The features.pkl file of this experiment is much larger compared to the other example. What is the reason of the different shape of the features?



Jason Brownlee March 24, 2018 at 6:29 am #

REPLY ↩

Perhaps the VGG model is cut at different points? e.g. keeping the dense or discarding it and working with the CNN outputs. I don't recall, but a comparison of the code would make it clearer.



Ashish March 28, 2018 at 4:56 pm #

REPLY ↩

Thanks for the tutorial, Jason. Here you are using features from a pretrained cnn model. But I want to backpropagate the error through the cnn as well and hence jointly train the cnn and lstm. i.e. load the cnn through a pretrained model and then train it further. Can you suggest a way to do that? Thanks.



Jason Brownlee March 29, 2018 at 6:31 am #

REPLY ↩

Yes, you can load the CNN as part of your in memory model.



Vector April 7, 2018 at 3:09 pm #

REPLY ↩

Hello Jason,

Thanks for this great posts and the one in here (<https://machinelearningmastery.com/develop-a-deep-learning-caption-generation-model-in-python/>)!!

I was trying to run the example you posted in here. I was able to train using different learning rate and more epoch on your “average pooling” model and saved them in a .h5 file.

However when I try to use the code in (<https://machinelearningmastery.com/develop-a-deep-learning-caption-generation-model-in-python/>) to Generate New Captions with the “Average Model” I trained, I got this error:

ValueError: Error when checking : expected input_1 to have 4 dimensions, but got array with shape (1, 4096)

Just wonder do you know how to use this .h5 file i just trained to make prediction on one other photo? Thanks!!



Jason Brownlee April 8, 2018 at 6:14 am #

REPLY ↩

It looks like your photo features might have too many dimensions. Change your code to provide the photo pixels directly. e.g.: photo[0]



Md. Zakir Hossain May 25, 2018 at 6:46 pm #

REPLY ↩

Hai Jason,

that's a great post. Really it is very helpful for me. However, I get the following error:

File “”, line 1, in

runfile('C:/Users/33083707/Codes/Projects/Final.py', wdir='C:/Users/33083707/Codes/Projects')

File “C:/Users/33083707/Anaconda3/lib/site-packages/spyder/utils/site/sitecustomize.py”, line 866, in runfile
execfile(filename, namespace)

File "C:\Users\33083707\Anaconda3\lib\site-packages\spyder\utils\site\sitecustomize.py", line 102, in
execfile

exec(compile(f.read(), filename, 'exec'), namespace)

File "C:/Users/33083707/Codes/Projects/Final.py", line 244, in

test_score = evaluate_model(model, test_descriptions, test_features, tokenizer, max_length)

File "C:/Users/33083707/Codes/Projects/Final.py", line 202, in evaluate_model

bleu = corpus_bleu(actual, predicted)

File "C:\Users\33083707\Anaconda3\lib\site-packages\nltk\translate\bleu_score.py", line 146, in corpus_bleu

p_i = modified_precision(references, hypothesis, i)

File "C:\Users\33083707\Anaconda3\lib\site-packages\nltk\translate\bleu_score.py", line 287, in

modified_precision

return Fraction(numerator, denominator, _normalize=False)

File "C:\Users\33083707\Anaconda3\lib\fractions.py", line 186, in __new__

raise ZeroDivisionError('Fraction(%s, 0)' % numerator)

ZeroDivisionError: Fraction(0, 0)



Jason Brownlee May 26, 2018 at 5:51 am #

REPLY ↩

Sorry to hear that, I have some suggestions to try here:

<https://machinelearningmastery.com/faq/single-faq/why-does-the-code-in-the-tutorial-not-work-for-me>



Mhemmed Elly June 29, 2018 at 9:29 pm #

REPLY ↩

Hello Dr. Jason,

I want to try this on my own data, but I don't know how to prepare it. I have a folder of images and a CSV file. Each line in the csv file contains an image ID/name and a caption that describes the image. Any advice on how to proceed?

Thank you



Jason Brownlee June 30, 2018 at 6:07 am #

REPLY ↩

Start by writing code to load the images into memory.

Perhaps you can use PIL or Pillo to load the images?

Jeet Sen Sarma July 6, 2018 at 6:53 pm #

REPLY ↩



in the following part of the code :

```
# run experiment
train_results, test_results = list(), list()
for i in range(n_repeats):
    # define the model
    model = define_model(vocab_size, max_length)
    # fit model
    model.fit_generator(data_generator(train_descriptions, train_features, tokenizer, max_length,
    n_photos_per_update), steps_per_epoch=n_batches_per_epoch, epochs=n_epochs, verbose=verbose)
    # evaluate model on training data
    train_score = evaluate_model(model, train_descriptions, train_features, tokenizer, max_length)
    test_score = evaluate_model(model, test_descriptions, test_features, tokenizer, max_length)
    # store
    train_results.append(train_score)
    test_results.append(test_score)
    print('>%d: train=%f test=%f' % ((i+1), train_score, test_score))
```

why are you defining the model within the for loop ? shouldn't it be defined outside of it ?



Jason Brownlee July 7, 2018 at 6:13 am #

REPLY ↩

Nope, I want a new model (random weights) for each repeat.

Why would I define outside the loop?



Jeet July 9, 2018 at 2:47 pm #

REPLY ↩

Could you please explain why is that? I think that a model has to be defined, and the weights are to be trained by learning from examples. But defining a new model everytime would delete the previously learned weights. Where am I going wrong?



Jason Brownlee July 10, 2018 at 6:41 am #

REPLY ↩

Yes, we want to throw it away. We are repeating the experiment with a model each time so we can see how well the model on our data performs on average. Not the performance of any one randomly fit model.

You can learn more here:

<https://machinelearningmastery.com/evaluate-skill-deep-learning-models/>



Md. Zakir Hossain July 9, 2018 at 6:13 pm #

REPLY ↩

Hi Jason, Thank you very much.

I tried to use DenseNet-121 model instead of VGG16 model to extract the features. But I got an error like:
ValueError: Error when checking input: expected input_1 to have 4 dimensions, but got array with shape (21, 1024).

Can you please suggest me anything about this.



Jeet July 10, 2018 at 4:26 am #

REPLY ↩

This is because both the CNN models have different output dimensions of their last layer from where you are pulling the features for your images. Best guess is to convert the output dimension of your DenseNet model into the same dimension of that of VGG-16.
Also, if you use Flatten() instead of GlobalMaxPooling2D() then I guess you will not face the same problem.



Jason Brownlee July 10, 2018 at 6:44 am #

REPLY ↩

Not really, I don't know about the things you are trying.



Omnia October 24, 2018 at 8:25 am #

REPLY ↩

Hi Jason

Thanks for the wonderful tutorial as always

I got this result on 70 epochs and 3 repeats
it seems that the accuracy increases when we have a large number of epochs

but I don't understand why I got the train and test = 0

this is my result

Epoch 70/70

– 14s – loss: 1.8873 – acc: 0.4100

Actual: startseq child and woman are at waters edge in big city endseq

Predicted: startseq child woman are edge edge in big endseq

Actual: startseq boy with stick kneeling in front of goalie net endseq

Predicted: startseq boy boy in in front and and of of of of of in front and of of in front boy front and of of

Actual: startseq woman crouches near three dogs in field endseq

Predicted: startseq two dog dogs in near near near near near near near near near near near near near near near

BLEU may give warnings when some predictions are shorter than expected. This might help understand the BLEU score calculation some more:

<https://machinelearningmastery.com/calculate-bleu-score-for-text-python/>

Leave a Reply

Name (required)

Email (will not be published) (required)

Website

SUBMIT COMMENT

Welcome to Machine Learning Mastery!




Hi, I'm Jason Brownlee, PhD

I write tutorials to help developers (*like you*) get results with machine learning.

[Read More](#)

Deep Learning for NLP

Develop deep learning models for text data.

[Click to Get Started Now!](#)

POPULAR

**How to Develop a Neural Machine Translation System from Scratch**

JANUARY 10, 2018

**So, You are Working on a Machine Learning Problem...**

APRIL 4, 2018

**How to Develop an N-gram Multichannel Convolutional Neural Network for Sentiment Analysis**

JANUARY 12, 2018

**How to Make Predictions with Keras**

APRIL 9, 2018

**11 Classical Time Series Forecasting Methods in Python (Cheat Sheet)**

AUGUST 6, 2018

**You're Doing it Wrong. Why Machine Learning Does Not Have to Be So Hard**

MARCH 28, 2018

**How to Develop LSTM Models for Multi-Step Time Series Forecasting of Household Power Consumption**

OCTOBER 10, 2018

You might also like...

- [How to Install Python for Machine Learning](#)

- [Your First Machine Learning Project in Python](#)
- [Your First Neural Network in Python](#)
- [Your First Classifier in Weka](#)
- [Your First Time Series Forecasting Project](#)

© 2019 Machine Learning Mastery. All Rights Reserved.

[Privacy](#) | [Disclaimer](#) | [Terms](#) | [Contact](#)